

## Exercise Round 1 Daniel Kusnetsoff

### Task 1

a) Convert the four points below (cartesian x,y-coordinates) into their corresponding homogeneous coordinate form.

$$x1 = [2; 1]$$

$$x1 = \begin{matrix} 2 \times 1 \\ 2 \\ 1 \end{matrix}$$

$$x1 = [2; 1; 1]$$

$$x1 = \begin{matrix} 3 \times 1 \\ 2 \\ 1 \\ 1 \end{matrix}$$

$$x2 = [1; -2]$$

$$x2 = \begin{matrix} 2 \times 1 \\ 1 \\ -2 \end{matrix}$$

$$x2 = [1; -2; 1]$$

$$x2 = \begin{matrix} 3 \times 1 \\ 1 \\ -2 \\ 1 \end{matrix}$$

$$x3 = [1; 1]$$

$$x3 = \begin{matrix} 2 \times 1 \\ 1 \\ 1 \end{matrix}$$

$$x3 = [1; 1; 1]$$

$$x3 = \begin{matrix} 3 \times 1 \\ 1 \\ 1 \\ 1 \end{matrix}$$

$$x4 = [-1; 0]$$

$$x4 = \begin{matrix} 2 \times 1 \\ -1 \\ 0 \end{matrix}$$

$$x4 = [-1; 0; 1]$$

$$\mathbf{x4} = \begin{pmatrix} 3 \\ -1 \\ 0 \\ 1 \end{pmatrix}$$

b) The line  $\mathbf{l}$  through two points  $\mathbf{x}$  and  $\mathbf{x}'$  is  $\mathbf{l} = \mathbf{x} \times \mathbf{x}'$ . Use this to form two lines, line  $\mathbf{l1}$  through homogeneous points  $\mathbf{x1}$  and  $\mathbf{x2}$ , and  $\mathbf{l2}$  through  $\mathbf{x3}$  and  $\mathbf{x4}$ .

As  $\mathbf{l} = \mathbf{x} \times \mathbf{x}'$ ,

$$\mathbf{x}(\text{transpose}) * \mathbf{l} = \mathbf{x}(\text{transpose}) * \mathbf{x} \text{ cross}(\mathbf{x}') = 0$$

$$\mathbf{A} = \text{cross}(\mathbf{x1}, \mathbf{x2})$$

$$\mathbf{A} = \begin{pmatrix} 3 \\ 3 \\ -1 \\ -5 \end{pmatrix}$$

$$\mathbf{B} = \text{cross}(\mathbf{x3}, \mathbf{x4})$$

$$\mathbf{B} = \begin{pmatrix} 3 \\ 1 \\ -2 \\ 1 \end{pmatrix}$$

c. The intersection of two lines  $\mathbf{l}$  and  $\mathbf{l}'$  is the point  $\mathbf{x} = \mathbf{l} \times \mathbf{l}'$ . Use lines  $\mathbf{l1}$  and  $\mathbf{l2}$  to calculate their point of intersection and convert this back into cartesian coordinates.

$$\mathbf{C} = \text{cross}(\mathbf{A}, \mathbf{B})$$

$$\mathbf{C} = \begin{pmatrix} 3 \\ -11 \\ -8 \\ -5 \end{pmatrix}$$

$$-11/5$$

$$\text{ans} = -2.2000$$

$$-8/5$$

$$\text{ans} = -1.6000$$

Point of intersection [-2.2;-1.6]

# Untitled17

January 16, 2022

```
[ ]: # -*- coding: utf-8 -*-
      """
      Created on Wed Jan 13 08:38:35 2021

      @author: tiitu
      """

import os
import sys
sys.path.append(os.getcwd())

from matplotlib.pyplot import imread
from skimage.transform import resize as imresize
#from scipy.misc import imresize # deprecated, may work with older versions of
↳scipy
import numpy as np
import matplotlib.pyplot as plt
from scipy.ndimage.filters import convolve as conv2
from scipy.ndimage.filters import convolve1d as conv1
from utils import imnoise, gaussian2, bilateral_filter

# Load test images and convert to double precision in the interval [0,1].
im = imread('einsteinpic.jpg') / 255.
im = imresize(im, (256, 256))

# Generate noise
imns = imnoise(im, 'salt & pepper', 0.05) * 1. # "salt and pepper"
↳noise
imng = im + 0.05*np.random.randn(im.shape[0],im.shape[1]) # zero-mean Gaussian
↳noise

# Apply a Gaussian filter with a standard deviation of 2.5
sigmad = 2.5
g, _, _, _, _, _ = gaussian2(sigmad)
```

```

gflt_imns = conv2(imns, g, mode='reflect')
gflt_imng = conv2(imng, g, mode='reflect')

# Instead of directly filtering with g, make a separable implementation
# where you use horizontal and vertical 1D convolutions.
# Store the results again to gflt_imns and gflt_imng, use conv1 instead.
# The result should not change.
# See Szeliski's Book chapter 3.2.1 Separable filtering, numpy.linalg.svd and
↳ scipy.ndimage.filters.convolve1d

##--your-code-starts-here--##
## 1d-gaussian

def gaussian1(sigma, N=None):
    if N is None:
        N = 2 * np.maximum(4, np.ceil(6*sigma)) + 1
    k = (N - 1) / 2.
    x = np.arange(-k, k+1)
    g = 1/(np.sqrt(2 * np.pi * sigma**2)) * np.exp(-(x**2) / (2 * sigma ** 2))
    return g

g1d=gaussian1(sigmad)

gflt_imns_x = conv1(imns, g1d, mode="reflect", axis=0)
gflt_imns_xy = conv1(gflt_imns_x, g1d, mode='reflect', axis=1)

gflt_imns_y = conv1(imns, g1d, mode="reflect", axis=1)
gflt_imns_yx = conv1(gflt_imns_y, g1d, mode='reflect', axis=0)

#gflt_imns = conv1(imns, g, mode='reflect')
#gflt_imng = conv1(imng, g, mode='reflect')
##--your-code-ends-here--##

# Median filtering is done by extracting a local patch from the input image
# and calculating its median
def median_filter(img, wsize):
    nrows, ncols = img.shape
    output = np.zeros([nrows, ncols])
    k = (wsize - 1) / 2

    for i in range(nrows):
        for j in range(ncols):
            # Calculate local region limits
            iMin = int(max(i - k, 0))
            iMax = int(min(i + k, nrows - 1))

```

```

jMin = int(max(j - k, 0))
jMax = int(min(j + k, ncols - 1))

# Use the region limits to extract a patch from the image,
# calculate the median value (e.g using numpy) from the extracted
# local region and store it to output using correct indexing.

##--your-code-starts-here--##

##--your-code-ends-here--##

return output

# Apply median filtering, use neighborhood size 5x5
# Store the results in medflt_imns and medflt_imng
# Use the median_filter function above

##--your-code-starts-here--##

##--your-code-ends-here--##

# Apply bilateral filter to each image with window size 11.
# See section 3.3.1 of Szeliski's book
# Use sigma value 2.5 for the domain kernel and 0.1 for range kernel.

wsize = 11
sigma_d = 2.5
sigma_r = 0.1

bflt_imns = bilateral_filter(imns, wsize, sigma_d, sigma_r)
bflt_imng = bilateral_filter(imng, wsize, sigma_d, sigma_r)


# Display filtering results
fig, axes = plt.subplots(nrows=2, ncols=4, figsize=(16,8))
ax = axes.ravel()
ax[0].imshow(imns, cmap='gray')
ax[0].set_title("Noisy input image")
ax[1].imshow(gflt_imns, cmap='gray')
ax[1].set_title("Result of gaussian filtering")
#ax[2].imshow(medflt_imns, cmap='gray')
#ax[2].set_title("Result of median filtering")
ax[3].imshow(bflt_imns, cmap='gray')
ax[3].set_title("Result of bilateral filtering")
ax[4].imshow(imng, cmap='gray')
ax[5].imshow(gflt_imng, cmap='gray')
#ax[6].imshow(medflt_imng, cmap='gray')

```



```
ax[7].imshow(bflt_img, cmap='gray')  
plt.suptitle("Filtering results", fontsize=20)  
plt.show()
```

# Filtering results

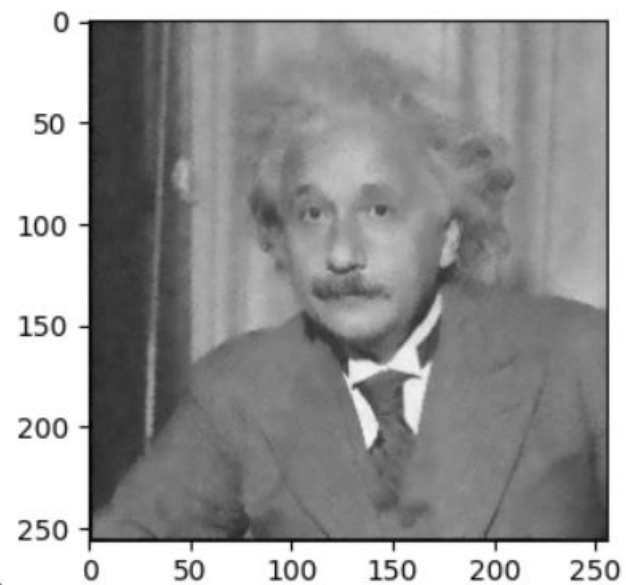



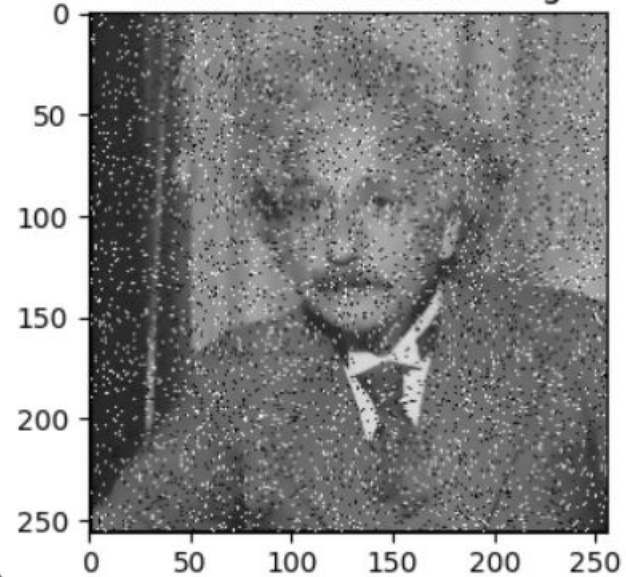
Noisy input image



Result of gaussian filtering



Result of bilateral filtering



# Untitled18

January 16, 2022

```
[ ]: import os
import sys
sys.path.append(os.getcwd())

from matplotlib.pyplot import imread
import numpy as np
from numpy.fft import fftshift, fft2
import matplotlib.pyplot as plt
from scipy.ndimage import gaussian_filter, map_coordinates
from utils import affinefit

# Load test images
man = imread('man.jpg') / 255.
wolf = imread('wolf.jpg') / 255.

# The pixel coordinates of eyes and chin have been manually found
# from both images in order to perform affine alignment
man_eyes_chin = np.array([[502, 465],    # left eye
                          [714, 485],    # right eye
                          [594, 875]])    # chin
wolf_eyes_chin = np.array([[851, 919],    # left eye
                           [1159, 947],    # right eye
                           [975, 1451]])    # chin

# Warp wolf to man using an affine transformation and the coordinates above
A, b = affinefit(man_eyes_chin, wolf_eyes_chin)
xv, yv = np.meshgrid(np.arange(0, man.shape[1]), np.arange(0, man.shape[0]))
pt = np.dot(A, np.vstack([xv.flatten(), yv.flatten()])) + np.tile(b, (xv.
    ↳size,1)).T
wolft = map_coordinates(wolf, (pt[1,:].reshape(man.shape), pt[0,:].reshape(man.
    ↳shape)))

# We'll start by simply blending the aligned images using additive
    ↳superimposition
additive_superimposition = man + wolft
```



```

# Next we create two different Gaussian kernels for low-pass filtering the two
→ images
sigmaA = 16
sigmaB = 8
man_lowpass = gaussian_filter(man, sigmaA, mode='nearest')
wolft_lowpass = gaussian_filter(wolft, sigmaB, mode='nearest')

# Your task is to create a hybrid image by combining a low-pass filtered
# version of the human face with a high-pass filtered wolf face
# HINT: A high-passed image is equal to the low-pass filtered result removed
→ from the original.
# Experiment also by trying different values for 'sigmaA' and 'sigmaB' above.

# Replace the zero image below with a high-pass filtered version of 'wolft'
##--your-code-starts-here--##
#wolft_highpass = np.zeros(wolft.shape)
wolft_highpass = wolft - wolft_lowpass
#plt(wolft_highpass)

#plt.show()
##--your-code-ends-here--##

# Replace also the zero image below with the correct hybrid image using your
→ filtered results
##--your-code-starts-here--##
#hybrid_image = np.zeros(man_lowpass.shape)
hybrid_image = man_lowpass + wolft_highpass
##--your-code-ends-here--##

# Try looking at the results from different distances.
# Notice how strongly the interpretation of the hybrid image is affected
# by the viewing distance
plt.figure(1)
plt.imshow(hybrid_image, cmap='gray')

# Display input images and both output images.
plt.figure(2)
plt.subplot(2,2,1)
plt.imshow(man, cmap='gray')
plt.title("Input Image A")
plt.subplot(2,2,2)
plt.imshow(wolft, cmap='gray')
plt.title("Input Image B")
plt.subplot(2,2,3)
plt.imshow(additive_superimposition, cmap='gray')
plt.title("Additive Superimposition")

```

```

plt.subplot(2,2,4)
plt.imshow(hybrid_image, cmap='gray')
plt.title("Hybrid Image")

# Visualize the log magnitudes of the Fourier transforms of the original images.
# Your task is to calculate 2D fourier transform for wolf/man and their
    ↳ filtered results using fft2 and fftshift
##--your-code-starts-here--##
#F_man = np.zeros(man.shape)
#F_man_lowpass = np.zeros(man_lowpass.shape)
#F_wolft = np.zeros(wolft.shape)
#F_wolft_highpass = np.zeros(wolft_highpass.shape)
from numpy import fft
## magnitudes
def shift_tf(image):
    return fft.fftshift(fft.fft2(image))


F_man = shift_tf(man)
F_man_lowpass = shift_tf(man_lowpass)
F_wolft = shift_tf(wolft)
F_wolft_highpass = shift_tf(wolft_highpass)
##--your-code-ends-here--##

# Display the Fourier transform results
plt.figure(3)
plt.subplot(2,2,1)
plt.imshow(np.log(np.abs(F_man)), cmap='gray')
plt.title("log(abs(F_man))")
plt.subplot(2,2,2)
plt.imshow(np.log(np.abs(F_man_lowpass)), cmap='gray')
plt.title("log(abs(F_man_lowpass)) image")
plt.subplot(2,2,3)
plt.imshow(np.log(np.abs(F_wolft)), cmap='gray')
plt.title("log(abs(F_wolft)) image")
plt.subplot(2,2,4)
plt.imshow(np.log(np.abs(F_wolft_highpass)), cmap='gray')
plt.title("log(abs(F_wolft_highpass))")


plt.show()

```


Input Image A




Input Image B

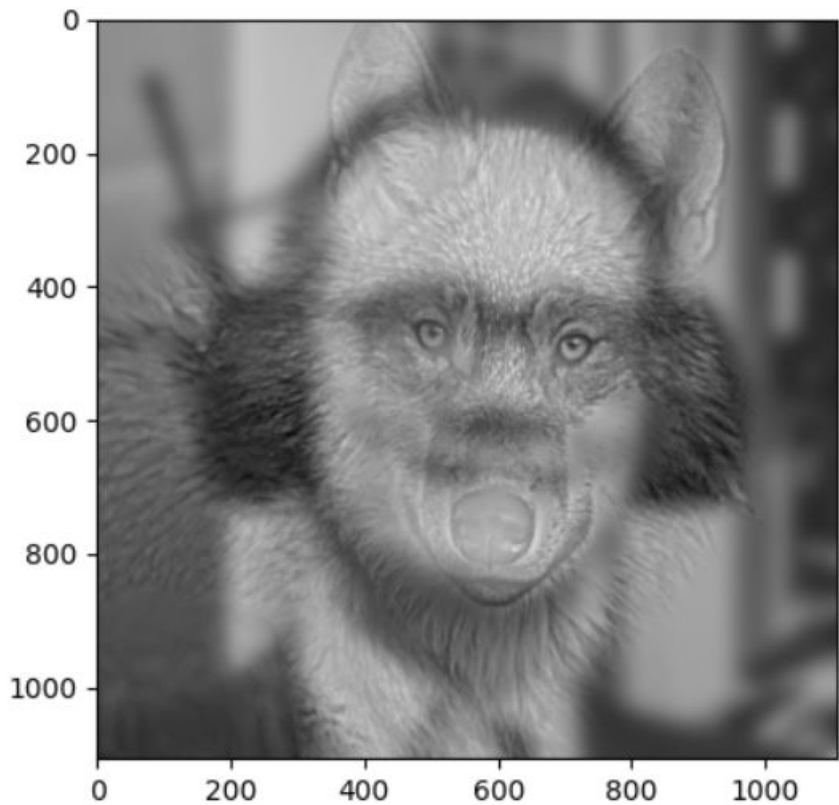


Additive Superimposition




Hybrid Image







$\log(\text{abs}(F_{\text{man}}))$



$\log(\text{abs}(F_{\text{man\_lowpass}}))$  image



$\log(\text{abs}(F_{\text{wolft}}))$  image



$\log(\text{abs}(F_{\text{wolft\_highpass}}))$

